Test a deployed agent with AgentGate

1. Deploy the agent/API and make sure its endpoint is reachable from the machine or CI runner where AgentGate will run. Use the deployed URL—not `localhost` unless both are running in the same environment.

2. Configure the API URL in `agent-search.yaml`:
   ```yaml
   http:
     url: https://your-deployed-host/api/search/parse
   ```
   Check that the request template and response mapping match the deployed API. The current config sends `query` and `variables` and reads `specialtyId`. Add required authentication using environment variables or CI secrets; don't commit credentials.

3. Include the AgentGate test code and dataset in the environment where the test will run. Ensure Python and AgentGate dependencies are installed. The dataset is `golden_dataset.csv`; it contains 99 queries with expected specialty IDs.

4. Run the smoke test from `Agent_testing\agentgate`:
   ```powershell
   python -m agentgate run --tier smoke --agent config\agent-search.yaml --dataset dataset\golden_dataset.csv
   ```

   Or run the launcher from the repository root:
   ```powershell
   .\Agent_testing\agentgate\run_multiple_queries.ps1
   ```

5. Review the result. AgentGate sends one request per dataset row. Reports are generated in `reports\<run-id>\`, including `report.xlsx`, `report.pdf`, and `results.json`; `reports\latest\` contains the latest copies. A nonzero exit code indicates the configured gate failed, for example because too many responses did not match expectations.

6. Automate after deployment. In your CI/CD pipeline, check out the test code, install dependencies, provide the deployed API URL and any credentials securely, wait for the API to become ready, then run the same command. Save the `reports` folder as a CI artifact for review.