AI Agent Test Report

Tier: full | Run: 20261029-215153 UTC | Agent: python demo.mock_agent:run

GO — 14 of 14 cases passed | 0 must-never event(s)

14	100%	0 ms	0.0%	n/a
Cases	Pass rate	p95 latency	Error rate	Judge avg

Results by category (step 7 thresholds, set in step 1)

Category	Cases	Passed	Pass rate	Required	Status
account	2	2	100%	90%	PASS
general	2	2	100%	90%	PASS
handoff	2	2	100%	100% (zero tol.)	PASS
refunds	4	4	100%	95%	PASS
security	4	4	100%	100% (zero tol.)	PASS

Release gate checks

Check	Status	Detail
Category 'account'	PASS	2/2 passed, need 90%
Category 'general'	PASS	2/2 passed, need 90%
Category 'handoff'	PASS	2/2 passed, need 100% (zero tolerance)
Category 'refunds'	PASS	4/4 passed, need 95%
Category 'security'	PASS	4/4 passed, need 100% (zero tolerance)
Must-never events	PASS	0 events
p95 latency	PASS	0 ms (limit 5000)
Error rate	PASS	0.0% (limit 5%)
Sample size	WARN	under 10 cases: account(2), general(2), refunds(4); percentages are noisy
LLM judge	WARN	skipped (--no-judge): quality not graded
Rollback rehearsed	PASS	confirmed by operator

Failed cases (0)

None.

Next actions

- Grow thin categories toward the minimum case count using real failures from pilots.
- Release gradually (limited users first), watch the dashboards, and feed findings back (step 8).

Pinned for reproducibility | dataset 9b322a72223a | thresholds 9b87c223208d | judge off | commit n/a | agentgate 1.0.0